

Attention and Deep Work

by *Diego Rincón* · *phronesis.world*

Your attention has an author. This course is about making sure it's you.

Who this is for

You reach for your phone without deciding to. You sit down to work and surface forty minutes later having done something else. You read less than you used to, think less than you used to, and you can feel it — a thinness in the days, a low hum of unease you can't name. You've tried deleting apps. It came back.

This course is not a productivity system. It's a map of one specific displacement — the gap between the attention you'd author and the attention that gets harvested from you — and a tested route back. If your attention is already yours, you don't need this. If you're not sure, that uncertainty is the first data point.

One note before we start. Some of what's described here — restlessness, low mood, compulsive checking — overlaps with real clinical conditions. This is a framework for thinking, not medical advice — work with a clinician on real conditions.

Module 1: Two States of Attention

S0: attention you author

Start with the framework's first object: the ground state, **S0**. S0 is not perfection. It's not a monk's mind or eight hours of unbroken flow. S0 is simply this: at any given moment, the thing you are attending to is the thing you chose to attend to.

In S0 you might still scroll. You might still watch something dumb. The difference is authorship. You decided to do it, you know roughly when you'll stop, and stopping doesn't feel like tearing something. S0 attention has a shape: it goes somewhere, stays a while, and comes back when called.

Notice what S0 is not. It's not “no phone.” It's not “always working.” A person in S0 can spend an evening on the couch watching a film, fully displaced from productivity, fully aligned with intention. The variable is not the content. The variable is who chose.

S*: attention that gets captured

S* is your actual state — where you actually are, not where you'd choose to be. For most people now, S* looks like this: attention that moves in small, jerky arcs, none of them chosen. Check the phone. Check it again. Open a tab. Forget why. A notification arrives and you obey it like a summons. The day passes in fragments and at the end of it you cannot say what you attended to, because nothing was attended to. Things merely passed in front of you.

The defining feature of S* is that an external system is doing the choosing. Your attention still moves — it's not destroyed — but its movements are scheduled by software whose explicit design goal is that your attention stays on it. You are not the author. You're the page.

Displacement: the gap is the object

The distance between S_0 and S^* is the **displacement** — call it **xi**. The whole course operates on this one quantity.

Xi is not a vibe. You can estimate it. Take an honest day and ask: of my waking attention, how much went where I sent it, and how much went where it was pulled? You don't need a precise number. You need the honesty of the estimate. Most people, doing this for the first time, are quietly shocked. The pulled fraction isn't a slice of the day. It's the day, with slices of authorship cut into it.

Two more quantities complete the toolkit, and we'll use them throughout. **D(xi)** is the daily cost of staying displaced — what the gap charges you every single day you don't close it. **C_return** is the one-time cost of going back — the discomfort, the setup work, the withdrawal. The entire argument of this course is an inequality: **D(xi)** compounds and **C_return** doesn't, so almost any displacement worth noticing is worth returning from. We'll earn that claim module by module.

Practice

This week: Run a one-day attention audit. Pick a normal weekday. Carry a card or a note file. Every time you notice your attention has moved — to the phone, a tab, a notification — make a tally and mark it A (I chose this) or P (I was pulled). No judgment, no correction. Just count. At the end of the day, look at the ratio.

Reflect:

1. What was your A-to-P ratio, roughly? Did the number surprise you?
2. Describe your S_0 in one sentence — not an idealized self, just: where would your attention go if nothing were pulling it?
3. When today did you feel most like the author? What were the conditions?

Module 2: The Capture Machine

It's not weakness, it's engineering

Before the mechanics, settle one thing: the systems pulling your attention were built by teams of capable people whose job — measured, reviewed, bonused — is to increase the time you spend looking. You're not weak. You're outnumbered. Reframing this matters because shame is itself a capture mechanism: feeling bad about your phone use is a feeling, and feelings get soothed by phones. The machine profits from your guilt about the machine.

So drop the morality. This is mechanics. Three mechanisms do most of the work.

Variable reward

A reward that arrives every time becomes boring. A reward that arrives unpredictably becomes compulsive. This is old, well-established behavioral science — it's why slot machines work — and the feed is a slot machine where the lever is your thumb.

Pull: nothing interesting. Pull: nothing. Pull: something genuinely delightful. That schedule — most pulls empty, occasional pulls rewarding, no way to predict which — trains the checking behavior harder than any consistent reward could. The emptiness of most of your scrolling isn't a

bug in the system. It's the payload. If the feed were reliably good, you'd check it like you check a weather forecast: once, on purpose. It's reliably *uncertain*, so you check it like a gambler.

Watch yourself open an app, find nothing, close it, and open it again within a minute. That loop is not information-seeking. It's lever-pulling.

False urgency

The second mechanism dresses the trivial as the urgent. The red badge, the buzz, the banner — these are interrupt signals, a visual language borrowed from genuine emergencies, applied to someone liking a photo. Your nervous system can't tell the difference at the speed the decision gets made. The interrupt fires first; the evaluation of whether it mattered comes after — after you've already looked, which is the only outcome the system needed.

Run the test: of the last twenty notifications you received, how many needed a response within the hour? Within the day? At all? For most people the honest answer hollows out the entire category. Nearly everything arriving with the costume of urgency is, on inspection, a message that could have waited until you chose to look.

Switching costs

The third mechanism is the one that does the deep damage, and it's the least visible. When your attention switches contexts, it doesn't switch clean. Part of your mind stays with the previous thing — researchers call this attention residue — and the cost of reassembling full focus on the original task is paid in minutes, not seconds.

Here's a toy model, numbers invented for shape, not fact. Suppose each interruption costs you ten minutes of reassembly before you're fully back. Suppose you're interrupted — or you interrupt yourself — four times an hour. That's forty minutes of every hour spent reassembling, which means you never reassemble. You spend the whole day in the shallow startup phase of thinking and never once reach the deep part. The work still gets done, sort of. But it's the work a much less capable person would do, because at four interruptions an hour, that's who you are.

This is why "I only check for a second" is the most expensive sentence in modern work. The check is a second. The reassembly is not. And the deep states — the ones where hard problems crack — require continuity that a single glance destroys.

Practice

This week: Disarm false urgency. Go into your phone's settings and turn off every notification that is not a live human being trying to reach you directly (calls, direct messages from actual people). Not "deliver quietly" — off. Leave it off for seven days before evaluating.

Reflect:

1. Which app do you check most compulsively, and what is its actual hit rate — out of ten checks, how many gave you something you were glad to see?
2. What did you feel in the first two days of silence? Name the feeling precisely.
3. What, honestly, would you miss if you saw it three hours late? Make the list. Look at how short it is.

Module 3: The Daily Bill

D(xi) is charged whether or not you look at it

Displacement has a daily cost, **D(xi)**, and its defining property is that it's charged invisibly. No invoice arrives. The cost shows up as absence — work not deepened, books not read, people not fully met — and absence doesn't trigger alarms. You don't feel the loss of a thought you never got to have.

So this module makes the bill visible. Four line items.

Shallow work

The switching-cost mechanics from Module 2 have a career-level consequence: there is a kind of work you can no longer do. Not “do worse” — *cannot do*, because it only exists past a threshold of continuity you no longer cross. Real writing, real design, real analysis, real strategy: these begin somewhere past the twenty- or thirty-minute mark of unbroken focus, in territory the fragmented day never reaches.

The cruelty of this cost is that shallow work still looks like work. The day is full. Email is answered, messages are handled, small tasks close. You're exhausted by evening, which feels like evidence of effort. But the substantive thing — the thing that would have justified the year — never happens, because it needed a depth of attention that was sold off in thirty-second increments. Busyness is what displacement looks like from the inside.

Lost learning

Learning anything real — a language, a skill, a body of knowledge — requires effortful focus sustained over months. It is precisely the activity that capture destroys, because it's effortful now and rewarding later, and the machine in your pocket is effortless now and empty later. Every time the hard-now/good-later activity competes head-to-head against the easy-now/empty-later one, easy-now wins the moment.

Notice the long arithmetic, though — toy model again. Suppose displaced attention eats one hour a day that S0-you would have spent learning. An hour a day, over a few years, is roughly the time cost of genuine competence in something — an instrument, a language, a field. That competence is gone. Not postponed. Gone, in the sense that the years went by and it doesn't exist. **D(xi)** compounds; that's the entire menace of it. The daily charge looks small. The decade total is a different person you didn't become.

Absent presence

The third line item is paid by other people. You can be in the room and not in the room. Children clock this with brutal accuracy: they know, before they can articulate it, whether your attention is on them or merely near them. Partners know. Friends know. The phone face-down on the dinner table isn't neutral — its presence alone splits your attention, because part of you is tracking it.

Connection runs on attention the way conversation runs on listening. Drain the attention and the forms of connection persist — dinners happen, conversations happen — but thinner, and everyone present can feel the thinness without being able to say its name. This cost never appears in any personal-productivity accounting, and it's the one people regret most.

Ambient anxiety

The fourth line item you carry in your body. A mind trained to check develops a hum — a low-grade, persistent vigilance, the feeling that something somewhere needs looking at. It doesn't switch off when the phone is away, because by then it's not the phone, it's the training. Boredom has become intolerable. Stillness feels itchy. The first thirty seconds of any queue, any red light, any pause produces an automatic reach.

This matters beyond comfort. The resting state of the mind — the unfocused drift where nothing is being consumed — is where consolidation happens, where the day's material composts into insight. A mind that's never bored is a mind that never digests. The hum costs you twice: the anxiety itself, and the thinking that the silence would have produced.

Practice

This week: Price your own D(xi). Take your audit from Module 1 and write the bill in words, not numbers — one sentence per line item. *The work I am not doing is .* ***The thing I am not learning is .*** *The person not getting my full attention is .* ***The hum shows up for me as .*** Put the four sentences somewhere you'll see them.

Reflect:

1. Which of the four costs is largest for you? Which one stung to write?
2. Project the toy arithmetic forward honestly: if the next five years look like the last six months, what doesn't get built?
3. Who in your life would notice first if your presence came back fully? What would they notice?

Module 4: The Return

C_return is real, finite, and front-loaded

Going back to S0 has a cost — **C_return** — and pretending otherwise is how most attempts fail. The return costs you setup work, social friction, and a stretch of genuine discomfort — often a week or two, though the shape varies by person. But hold the two costs side by side: C_return is paid once and ends. D(xi) is paid daily and compounds. Any finite one-time cost beats an unbounded recurring one. The math isn't close. The only reason the decision feels close is that C_return is paid up front and in concentrated form, while D(xi) is diffuse and invisible. You're comparing a visible cost to an invisible one, and the visible one looms.

The protocol below spreads C_return over four weeks, in deliberate order: environment first, then tolerance, then capacity. Don't reorder it. Each stage makes the next one affordable.

Week one: environment before willpower

Rule one of the return: never use willpower for what architecture can do. Willpower is a small, depletable budget, and spending it on resisting a slot machine that's within arm's reach is spending it at the worst possible exchange rate. So week one touches no habits at all. It only moves objects.

Do these, mechanically, like setting up a workshop:

- **Notifications:** already off from Module 2. If you skipped it, it's now mandatory.
- **The phone gets a home** that is not your pocket and not your desk — a drawer, a shelf in another room. It lives there by default and you visit it on purpose.

- **The first hour of the day is offline.** Not as discipline — as architecture. The phone charges outside the bedroom; you don't see a screen until you've been awake an hour. The first input of the day sets the register of the day.
- **Remove the slot machines from the phone.** The feeds you pull most compulsively: delete the app. The web version, with a login wall in the way, is fine — friction is the point. You're not banning content; you're banning the one-thumb lever.
- **One screen at work.** Whatever you're working on is the only thing visibly open. Every additional visible surface is a standing invitation to switch.

That's the week. Nothing heroic. You're not fighting the machine yet — you're moving it across the room so that fighting becomes mostly unnecessary.

Week two: retraining boredom

With the environment quieted, the second week works on you. The capacity you're rebuilding is tolerance for unstimulated time — boredom, the load-bearing skill underneath all deep attention. A mind that can't be bored can't be focused, because focus *is* the willingness to stay with one thing after its novelty is gone.

The retraining is almost embarrassingly simple: put gaps back in the day and don't fill them. Wait in the line without reaching. Walk without audio. Eat one meal a day with no input at all. Sit with your coffee and just sit. Each gap will itch, and the itch is the rep — every time you feel the pull and don't act on it, the wiring weakens slightly. You're not doing nothing in those gaps. You're doing the work.

Keep the gaps small and frequent rather than long and rare. Thirty unstimulated moments in a week beats one heroic silent Sunday, for reasons Module 5 will make formal.

Weeks three and four: deep blocks

Now you have a quiet environment and a mind that can sit still. Spend it. The deep block is the payoff stage: a protected stretch of time, one important thing, zero inputs.

The form: pick the work that matters most — the thing from your Module 3 bill, the work you weren't doing. Schedule a block at the same time each day, because regularity removes the daily negotiation. Start at a length you can actually complete with full integrity — better a shorter block kept perfectly than a longer one kept approximately, because what you're training is the inviolability, not the duration. Phone in its home, one screen, door closed if you have a door. When you drift — you will — the rep is the noticing and the return, without commentary. Drift, notice, return. That loop is the whole skill. Extend the block gradually as the returns get easier.

Two blocks a day is plenty. Deep attention is expensive tissue; it doesn't run all day, and pretending it does just turns deep blocks into shallow blocks with better branding. Protect a small amount absolutely rather than a large amount approximately.

Withdrawal: what it feels like, and that it passes

Be told in advance so it doesn't ambush you. The early return feels bad in specific ways: restlessness with a physical edge; phantom reaches for the phone; a strange spike of anxiety in quiet moments, as if the silence itself were a problem; for some people, irritability and a flat, grey mood as the

brain recalibrates away from a constant drip of cheap salience. You may feel a kind of grief, and underneath it the suspicion that without the noise you are boring — that there’s nothing in there.

There is. It’s just slower than what you’re used to. The noise wasn’t covering emptiness; it was preventing depth from getting a word in. Most people find the acute discomfort peaks early in the process and fades over a couple of weeks, somewhere in week two or three of this protocol — not as a guarantee, but as the common shape. What replaces it tends to arrive quietly: a day that feels longer in the good way, a thought that completes itself, a book that holds you again. If instead of fading the distress deepens, treat that seriously — this is a framework, not treatment, and a clinician is the right next call.

Practice

This week: Execute week one, fully — every environmental change on the list, in one sitting, today if possible. Don’t improve the list. Don’t negotiate with it. The whole stage takes under an hour of actual work.

Reflect:

1. Which environmental change did you resist most? What does that resistance tell you about which lever owns you?
2. What is the one piece of work your deep blocks will be for? Name it specifically — “my book’s second chapter,” not “writing.”
3. What’s your honest prediction of what withdrawal will feel like for you? Write it down now, so future-you can check it.

Module 5: Staying Back

The maintenance theorem

Here is where most attention advice dies, so look at the failure mode directly. The popular model is the detox: live displaced for months, then take a dramatic abstention — a screen-free week, a retreat — and return purified. It fails predictably, and the framework says exactly why.

The **maintenance theorem**: returning on a schedule beats letting displacement accumulate, and the optimal return frequency balances the fixed cost of each return against the rate of drift. When drift is fast, returns must be frequent — that’s the whole theorem applied.

Attention drift is *fast*. The capture machinery doesn’t pause while you’re being virtuous; the moment maintenance stops, the pull resumes, and x_i starts growing again that same afternoon. Let it grow for three months and you’re paying three months of compounding $D(x_i)$, and the eventual return is a full C_{return} all over again — expensive, dreaded, postponed, which is exactly how three months becomes nine months. The detox model is the maximum-cost strategy dressed up as discipline: maximum accumulated daily cost, maximum repeated return cost, repeated forever.

The alternative: return so often that displacement never accumulates enough to charge you. A daily boundary is a tiny return — near-zero cost, because x_i never got large. Daily boundaries beat monthly detoxes not by a little but structurally, the way brushing your teeth beats annual dental surgery.

The daily returns

Maintenance is a short list of boundaries held every day. Yours will vary; a working set looks like this:

- **First hour offline, last hour offline.** The day's bookends belong to you. Capture owns neither your waking register nor your falling-asleep mind.
- **The phone keeps its home.** Permanently. The drawer is not a punishment; it's where the tool lives, the way the hammer lives in the toolbox and not in your hand.
- **One deep block, daily, non-negotiable.** Even short. Especially on the days it seems impossible — those are the days drift is fastest.
- **Gaps stay empty.** The line, the walk, the wait: still no filler. This is the boredom muscle's daily maintenance dose.
- **A daily glance at xi.** Ten seconds at day's end: was I the author today, mostly? Not a journal, not a score — a glance. Drift caught at day one costs nothing to correct. That glance is the cheapest tool in this entire course.

The weekly check and the substrate

Once a week, take five minutes and ask three questions: Where did capture get back in this week? (It will — a new app, a reactivated notification, a feed crept back; this is weather, not failure.) Did the deep blocks hold, and did they go to the work that matters? Has any boundary quietly eroded enough to need re-setup? Fix what you find. Done.

And tend the substrate, because attention is a biological function before it's a virtue, and no protocol survives a wrecked substrate. Sleep is the single largest input — a fragmented night produces a fragmented mind, and no boundary compensates for it. Movement helps; a daily walk is one of the cheapest cognitive interventions available, and it doubles as a gap. Eat in ways that hold you steady rather than spiking and crashing you — oats, beans, lentils, nuts, fruit, vegetables, whole grains; the boring-sounding foods are the steady-attention foods. Caffeine is a fine tool early in the day and an attention tax when it bleeds into your sleep. None of this is exotic. All of it is load-bearing.

When you relapse — and you will

Some week — illness, a crisis, a deadline, travel — the boundaries collapse and you'll find yourself back in the feed, hours deep, the hum back in your chest. This is not a verdict. It's drift, which is the normal physics of the situation, and the framework already has the tool: notice, estimate the new xi, pay the return cost, resume. The return is cheaper this time — the architecture is built, the path is known, the deep-block habit has grooves. What ruins people is not the relapse but the story about the relapse: *I've broken it, the streak is dead, I'm someone who can't do this*. There is no streak. There's only a ground state and a distance from it, and the distance is always, always walkable.

Practice

This week: Install the daily glance. Pick a fixed anchor — closing the laptop, brushing your teeth — and attach the ten-second question to it: *author or page today?* Do it every day for the week, and let it trigger one small correction whenever the answer is “page.”

Reflect:

1. Which daily boundary will be hardest for you to hold? What architectural change would make it easier than willpower?
2. What was your last relapse-and-shame spiral (in any domain), and how would it have gone differently with “drift, not verdict” as the frame?
3. What’s your realistic re-capture vector — the specific app, person, or situation most likely to reopen the breach? What’s the pre-decided response?

The Worksheet

Fill this in for yourself — not in your head, on the page. One row per domain of your attention. Be concrete in every cell: real activities, real costs in words, a real decision with a date. The act of writing S^* honestly is half the return.

Domain	My S_0 (what authored attention looks like here)	My S^* (where I actually am)	Daily cost $D(x_i)$ (what this charges me, in words)	Return cost	
				C_{return} (what going back will take)	Decision (boundary + start date)
Morning					
Work hours					
Evenings					
With the people I love					
Gaps and waits					
Learning					

When the table is full, read down the $D(x_i)$ column in one pass. That column is your life’s current invoice. Then read the Decision column. That’s the cancellation notice. Sign it by starting.

The Path

You now have the whole structure. A ground state: attention you author. An actual state: attention that’s been captured by machinery built for exactly that purpose. A gap that charges you daily — in shallow work, lost learning, absent presence, and a hum in the chest — and a return that costs real discomfort exactly once.

The return order: environment, then boredom, then depth. The maintenance order: daily glance, daily boundaries, weekly check. The theorem underneath all of it: frequent small returns beat rare dramatic ones, always, because drift never sleeps and compounding never pauses.

Nothing in this course works as knowledge. It works as practice — the audit done, the notifications silenced, the phone in its drawer, the block held, the gap left empty, the glance at day’s end. Each one is small. That’s not a weakness of the method; it is the method. The author of your attention is rebuilt in ten-second decisions, made daily, starting with the one you make when you put this down.

Put it down. The first gap starts now. Don’t fill it.